

# Cutlass SGEMM Analysis

---

## Review of GEMM Implementation

### GEMM general structure

- infrastructure
  1.  $A \rightarrow m \times k$  &  $B \rightarrow k \times n$  &  $C \rightarrow m \times n$
  2.  $C = \alpha * A \times B^T + \beta * C$
- solution
  1.  $m = M * bm$  &  $n = N * bn$  &  $k = K * bk$
  2. a grid =  $M * N$  blocks, each block calculate a  $bm * bn$  area
  3. for a block, each time load  $bm * bk$  data from A and  $bk * bn$  data from B to shared memory
  4. there are K rounds for a block to calculate the target block
  5.  $bm = X * rm$  &  $bn = Y * rn$
  6. there are  $X * Y$  threads, and each thread calculate a block of  $rm * rn$
- implement

```
// Device function to compute a thread block's accumulated matrix product
__device__ void block_matrix_product(int K_dim) {

    // Fragments used to store data fetched from SMEM
    value_t frag_a[ThreadItemsY];
    value_t frag_b[ThreadItemsX];

    // Accumulator storage
    accum_t accumulator[ThreadItemsX][ThreadItemsY];

    // GEMM Mainloop - iterates over the entire K dimension - not unrolled
    for (int kblock = 0; kblock < K_dim; kblock += BlockItemsK) {

        // Load A and B tiles from global memory and store to SMEM
        //
        // (not shown for brevity - see the CUTLASS source for more detail)
        ...

        __syncthreads();

        // warp tile structure - iterates over the Thread Block tile
        #pragma unroll
        for (int warp_k = 0; warp_k < BlockItemsK; warp_k += WarpItemsK) {


            // Fetch frag_a and frag_b from SMEM corresponding to k-index
            //
            // (not shown for brevity - see CUTLASS source for more detail)
            ...
        }
    }
}
```

```

// Thread tile structure - accumulate an outer product
#pragma unroll
for (int thread_x = 0; thread_x < ThreadItemsX; ++thread_x) {
    #pragma unroll
    for (int thread_y=0; thread_y < ThreadItemsY; ++thread_y) {
        accumulator[thread_x][thread_y] += frag_a[y]*frag_b[x];
    }
}

__syncthreads();
}
}

```



- now the above optimization can reach 80% cublas

## Advance Optimization


- transpose matrix A ( smemA ) and use LDS.128 ( vector read instruction )
- try LDGSTS ?
- prefetch & read/write into different buffer (double buffer)

we use read SM & write SM to represent the two shared memory

1. load the data to write SM and write reg

2. for each iteration, switch read SM & write SM to avoid the latency of accessing mem

```
for k in 256 big_loop:
    prefetch next loop data to write_SM
    // compute in read_SM
    for iter in 8 small_loop:
        prefetch next loop data to write_REG
        compute in read_REG
```



- implement

```
#define TILE_K 16
__shared__ float4 smemA[2][TILE_K * 128 / 4];
__shared__ float4 smemB[2][TILE_K * 128 / 4];
float4 c[8][2] = {{make_float4(0.f, 0.f, 0.f, 0.f)}};
float4 ldg_a_reg[2];
float4 ldg_b_reg[2];
float4 a_reg[2][2];
float4 b_reg[2][2];

// transfer first tile from global mem to shared mem
load_gmem_tile_to_reg(A, 0, ldg_a_reg);
load_gmem_tile_to_reg(B, 0, ldg_b_reg);

store_reg_to_smem_tile_transpose(ldg_a_reg, 0, smemA[0]);
store_reg_to_smem_tile(ldg_b_reg, 0, smemB[0]);
__syncthreads();

// load first tile from shared mem to register
load_smem_tile_to_reg(smemA[0], 0, a_reg[0]);
load_smem_tile_to_reg(smemB[0], 0, b_reg[0]);

int write_stage_idx = 1; //ping pong switch
do {
    i += TILE_K;
    // load next tile from global mem
    load_gmem_tile_to_reg(A, i, ldg_a_reg);
    load_gmem_tile_to_reg(B, i, ldg_b_reg);

    int load_stage_idx = write_stage_idx ^ 1;

#pragma unroll
```

```

        for(int j = 0; j < TILE_K - 1; ++j) {
            // load next tile from shared mem to register
            load_smem_tile_to_reg(smemA[load_stage_idx], j + 1, a_reg[(j +
1) % 2]);
            load_smem_tile_to_reg(smemB[load_stage_idx], j + 1, b_reg[(j +
1) % 2]);
            // compute matrix multiply accumulate 8x8
            mma8x8(a_reg[j % 2], b_reg[j % 2], c);
        }

        if(i < K) {
            // store next tile to shared mem
            store_reg_to_smem_tile_transpose(ldg_a_reg, 0,
smemA[write_stage_idx]);
            store_reg_to_smem_tile(ldg_b_reg, 0, smemB[write_stage_idx]);
            // use double buffer, only need one sync
            __syncthreads();
            // switch
            write_stage_idx ^= 1;
        }

        // load first tile from shared mem to register of next iter
        load_smem_tile_to_reg(smemA[load_stage_idx ^ 1], 0, a_reg[0]);
        load_smem_tile_to_reg(smemB[load_stage_idx ^ 1], 0, b_reg[0]);
        // compute last tile mma 8x8
        mma8x8(a_reg[1], b_reg[1], c);
    } while (i < K);

    store_c(c, C);

```

- reach 97.5% cublas

## surpass cutlass

- SASS (Shader Assembly)
  - register bank conflict
  - register reuse

## Cutlass GEMM Implementation

### code structure

- interface

```
cutlass::gemm::device::Gemm<A_type,A_save,B_type,B_save,C_type,C_save>
```

- select gemm shape

```
cutlass/include/cutlass/gemm/device/gemm.h
```

```
template <
```

```

    /// Element type for A matrix operand
    typename ElementA_,
    /// Layout type for A matrix operand
    typename LayoutA_,
    /// Element type for B matrix operand
    typename ElementB_,
    /// Layout type for B matrix operand
    typename LayoutB_,
    /// Element type for C and D matrix operands
    typename ElementC_,
    /// Layout type for C and D matrix operands
    typename LayoutC_,
    /// Element type for internal accumulation
    typename ElementAccumulator_ = ElementC_,
    /// Operator class tag
    typename OperatorClass_ = arch::OpClassSimt,
    /// Tag indicating architecture to tune for
    typename ArchTag_ = arch::Sm70,
    /// Threadblock-level tile size (concept: GemmShape)
    typename ThreadblockShape_ = typename DefaultGemmConfiguration<
        OperatorClass_, ArchTag_, ElementA_, ElementB_, ElementC_,
        ElementAccumulator_>::ThreadblockShape,
    /// Warp-level tile size (concept: GemmShape)
    typename WarpShape_ = typename DefaultGemmConfiguration<
        OperatorClass_, ArchTag_, ElementA_, ElementB_, ElementC_,
        ElementAccumulator_>::WarpShape,
    /// Instruction-level tile size (concept: GemmShape)
    typename InstructionShape_ = typename DefaultGemmConfiguration<
        OperatorClass_, ArchTag_, ElementA_, ElementB_, ElementC_,
        ElementAccumulator_>::InstructionShape,
    /// Epilogue output operator
    typename EpilogueOutputOp_ = typename DefaultGemmConfiguration<
        OperatorClass_, ArchTag_, ElementA_, ElementB_, ElementC_,
        ElementAccumulator_>::EpilogueOutputOp,
    /// Threadblock-level swizzling operator
    typename ThreadblockSwizzle_ =
        typename threadblock::GemmIdentityThreadblockSwizzle<>,
    /// Number of stages used in the pipelined mainloop
    int Stages =
        DefaultGemmConfiguration<OperatorClass_, ArchTag_, ElementA_,
ElementB_,
                                ElementC_, ElementAccumulator_>::kStages,
    /// Access granularity of A matrix in units of elements
    int AlignmentA =
        DefaultGemmConfiguration<OperatorClass_, ArchTag_, ElementA_,
ElementB_,
                                ElementC_,
ElementAccumulator_>::kAlignmentA,
    /// Access granularity of B matrix in units of elements
    int AlignmentB =
        DefaultGemmConfiguration<OperatorClass_, ArchTag_, ElementA_,
ElementB_,
                                ElementC_,
ElementAccumulator_>::kAlignmentB,
    /// If true, kernel supports split-K with serial reduction

```

```

bool SplitKSerial = false,
/// Operation performed by GEMM
typename Operator_ = typename DefaultGemmConfiguration<
    OperatorClass_, ArchTag_, ElementA_, ElementB_, ElementC_,
    ElementAccumulator_>::Operator,
/// Gather operand A by using an index array
bool GatherA = false,
/// Gather operand B by using an index array
bool GatherB = false,
/// Scatter result D by using an index array
bool ScatterD = false>
class Gemm ;

```

cutlass/include/cutlass/gemm/device/default\_gemm\_configuration.h

```

template <
    typename OperatorClass,
    typename ArchTag,
    typename ElementA,
    typename ElementB,
    typename ElementC,
    typename ElementAccumulator
>
struct DefaultGemmConfiguration;

```

Kstages: 2 KAlignmentA: 1 KAlignmentB: 1 KAlignmentC: 1

- kernel call

cutlass/include/cutlass/gemm/device/gemm.h

```

Status operator()(cudaStream_t stream = nullptr) {
    return run(stream);
}

Status run(cudaStream_t stream = nullptr) {

    ThreadblockSwizzle threadblock_swizzle;

    dim3 grid =
threadblock_swizzle.get_grid_shape(params_.grid_tiled_shape);
    dim3 block(GemmKernel::kThreadCount, 1, 1);

    cudaError_t result;

    int smem_size = int(sizeof(typename GemmKernel::SharedStorage));

    if (smem_size >= (48 << 10)) {
        result = cudaFuncSetAttribute(Kernel<GemmKernel>,
cudaFuncAttributeMaxDynamicSharedMemorySize,
smem_size);
    }
}

```

```

        if (result != cudaSuccess) {
            return Status::kErrorInternal;
        }
    }

    // kernel call
    cutlass::Kernel<GemmKernel><<<grid, block, smem_size, stream>>>
(params_);
    // kernel call

    result = cudaGetLastError();

    return result == cudaSuccess ? Status::kSuccess :
Status::kErrorInternal;
}

```

cutlass/include/cutlass/device\_kernel.h

```

template <typename Operator>
__global__
void Kernel(typename Operator::Params params) {
    // Dynamic shared memory base pointer
    extern __shared__ int SharedStorageBase[];

    // Declare pointer to dynamic shared memory.
    typename Operator::SharedStorage *shared_storage =
        reinterpret_cast<typename Operator::SharedStorage *>
(SharedStorageBase);

    Operator op;

    op(params, *shared_storage);
}

```

- DefaultGemm structure

cutlass/include/cutlass/gemm/kernel/default\_gemm.h

```

template <
    /// Element type for A matrix operand
    typename ElementA,
    /// Layout type for A matrix operand
    typename LayoutA,
    /// Access granularity of A matrix in units of elements
    int kAlignmentA,
    /// Element type for B matrix operand
    typename ElementB,
    /// Layout type for B matrix operand
    typename LayoutB,
    /// Access granularity of B matrix in units of elements
    int kAlignmentB,
    /// Element type for C and D matrix operands
    typename ElementC,
    /// Element type for internal accumulation

```

```

typename ElementAccumulator,
/// Threadblock-level tile size (concept: GemmShape)
typename ThreadblockShape,
/// warp-level tile size (concept: GemmShape)
typename warpShape,
/// warp-level tile size (concept: GemmShape)
typename InstructionShape,
/// Epilogue output operator
typename EpilogueOutputOp,
/// Threadblock-level swizzling operator
typename ThreadblockSwizzle,
/// If true, kernel is configured to support serial reduction in the
epilogue
bool SplitkSerial,
/// Operation performed by GEMM
typename Operator,
/// Use zfill or predicate for out-of-bound cp.async
SharedMemoryClearOption SharedMemoryClear,
/// Gather operand A by using an index array
bool GatherA,
/// Gather operand B by using an index array
bool GatherB,
/// Scatter result D by using an index array
bool ScatterD
>
struct DefaultGemm<
    ElementA, LayoutA, kAlignmentA,
    ElementB, LayoutB, kAlignmentB,
    ElementC, layout::RowMajor,
    ElementAccumulator,
    arch::OpClassTensorOp,
    arch::Sm75,
    ThreadblockShape,
    warpShape,
    InstructionShape,
    EpilogueOutputOp,
    ThreadblockSwizzle,
    2,
    SplitkSerial,
    Operator,
    SharedMemoryClear,
    GatherA,
    GatherB,
    ScatterD
> {

    /// Define the threadblock-scoped matrix multiply-accumulate
    using Mma = typename cutlass::gemm::threadblock::DefaultMma<
        ElementA,
        LayoutA,
        kAlignmentA,
        ElementB,
        LayoutB,
        kAlignmentB,
        ElementAccumulator,

```



```

    layout::RowMajor,
    arch::OpClassTensorOp,
    arch::Sm75,
    ThreadblockShape,
    WarpShape,
    InstructionShape,
    2,
    operator,
    false,
    SharedMemoryClear,
    GatherA,
    GatherB
>::ThreadblockMma;

static const int kPartitionsK = ThreadblockShape::kk / warpShape::kk;

/// Define the epilogue
using Epilogue = typename
cutlass::epilogue::threadblock::DefaultEpilogueTensorOp<
    ThreadblockShape,
    typename Mma::Operator,
    kPartitionsK,
    EpilogueOutputOp,
    EpilogueOutputOp::kCount,
    ScatterD
>::Epilogue;

/// Define the kernel-level GEMM operator.
using GemmKernel = kernel::Gemm<Mma, Epilogue, ThreadblockSwizzle,
    SplitKSerial>;
};

```

cutlass/include/cutlass/gemm/kernel/gemm.h

```

CUTLASS_DEVICE
void operator()(Params const &params, SharedStorage &shared_storage) {

    // Compute threadblock location
    ThreadblockSwizzle threadblock_swizzle;

    cutlass::gemm::GemmCoord threadblock_tile_offset =
        threadblock_swizzle.get_tile_offset(params.swizzle_log_tile);

    // Early exit if CTA is out of range
    if (params.grid_tiled_shape.m() <= threadblock_tile_offset.m() ||
        params.grid_tiled_shape.n() <= threadblock_tile_offset.n()) {

        return;
    }

    // Compute initial location in logical coordinates
    cutlass::MatrixCoord tb_offset_A{
        threadblock_tile_offset.m() * Mma::Shape::kM,
        threadblock_tile_offset.k() * params.gemm_k_size,
    };
};

```

```

cutlass::MatrixCoord tb_offset_B{
    threadblock_tile_offset.k() * params.gemm_k_size,
    threadblock_tile_offset.n() * Mma::Shape::kN
};

// Problem size is a function of threadblock index in the K dimension
int problem_size_k = min(
    params.problem_size.k(),
    (threadblock_tile_offset.k() + 1) * params.gemm_k_size);

// Compute threadblock-scoped matrix multiply-add
int gemm_k_iterations = (problem_size_k - tb_offset_A.column() +
Mma::Shape::kK - 1) / Mma::Shape::kK;

// Compute position within threadblock
int thread_idx = threadIdx.x;

// Construct iterators to A and B operands
typename Mma::IteratorA iterator_A(
    params.params_A,
    params.ref_A.data(),
    {params.problem_size.m(), problem_size_k},
    thread_idx,
    tb_offset_A,
    params.gather_A_indices);

typename Mma::IteratorB iterator_B(
    params.params_B,
    params.ref_B.data(),
    {problem_size_k, params.problem_size.n()},
    thread_idx,
    tb_offset_B,
    params.gather_B_indices);

// Broadcast the warp_id computed by lane 0 to ensure dependent code
// is compiled as warp-uniform.
int warp_idx = __shfl_sync(0xffffffff, threadIdx.x / 32, 0);
int lane_idx = threadIdx.x % 32;

//
// Main loop
//

// Construct thread-scoped matrix multiply
Mma mma(shared_storage.main_loop, thread_idx, warp_idx, lane_idx);

typename Mma::FragmentC accumulators;

accumulators.clear();

if (!ksplitkSerial || gemm_k_iterations > 0) {
    // Compute threadblock-scoped matrix multiply-add
    mma(gemm_k_iterations, accumulators, iterator_A, iterator_B,
accumulators);

```

```

}

//
// Epilogue
//

OutputOp output_op(params.output_op);

//
// Masked tile iterators constructed from members
//

threadblock_tile_offset =
    threadblock_swizzle.get_tile_offset(params.swizzle_log_tile);

//assume identity swizzle
MatrixCoord threadblock_offset(
    threadblock_tile_offset.m() * Mma::Shape::kM,
    threadblock_tile_offset.n() * Mma::Shape::kN
);

int block_idx = threadblock_tile_offset.m() +
threadblock_tile_offset.n() * params.grid_tiled_shape.m();

// Construct the semaphore.
Semaphore semaphore(params.semaphore + block_idx, thread_idx);

// If performing a reduction via split-K, fetch the initial
synchronization
if (kSplitKSerial && params.grid_tiled_shape.k() > 1) {

    // Fetch the synchronization lock initially but do not block.
    semaphore.fetch();

    // Indicate which position in a serial reduction the output operator
is currently updating
    output_op.set_k_partition(threadblock_tile_offset.k(),
params.grid_tiled_shape.k());
}

// Tile iterator loading from source tensor.
typename Epilogue::OutputTileIterator iterator_C(
    params.params_C,
    params.ref_C.data(),
    params.problem_size.mn(),
    thread_idx,
    threadblock_offset,
    params.scatter_D_indices
);

// Tile iterator writing to destination tensor.
typename Epilogue::OutputTileIterator iterator_D(
    params.params_D,
    params.ref_D.data(),
    params.problem_size.mn(),

```

```

        thread_idx,
        threadblock_offset,
        params.scatter_D_indices
    );

    Epilogue epilogue(
        shared_storage.epilogue,
        thread_idx,
        warp_idx,
        lane_idx);

    // wait on the semaphore - this latency may have been covered by
    iterator construction
    if (ksplitKSerial && params.grid_tiled_shape.k() > 1) {

        // For subsequent threadblocks, the source matrix is held in the 'D'
        tensor.
        if (threadblock_tile_offset.k()) {
            iterator_C = iterator_D;
        }

        semaphore.wait(threadblock_tile_offset.k());
    }

    // Execute the epilogue operator to update the destination tensor.
    epilogue(output_op, iterator_D, accumulators, iterator_C);

    //
    // Release the semaphore
    //

    if (ksplitKSerial && params.grid_tiled_shape.k() > 1) {

        int lock = 0;
        if (params.grid_tiled_shape.k() == threadblock_tile_offset.k() + 1) {

            // The final threadblock resets the semaphore for subsequent grids.
            lock = 0;
        }
        else {
            // Otherwise, the semaphore is incremented
            lock = threadblock_tile_offset.k() + 1;
        }

        semaphore.release(lock);
    }
}

```

cutlass/include/cutlass/gemm/threadblock/mma\_pipelined.h

```

CUTLASS_DEVICE
void operator()(
    int gemm_k_iterations,           ///< number of
    iterations of the mainloop

```

```

    FragmentC &accum,                                     ///< destination
    accumulator tile
    IteratorA iterator_A,                                 ///<< iterator over A
    operand in global memory
    IteratorB iterator_B,                                 ///<< iterator over B
    operand in global memory
    FragmentC const &src_accum,                           ///<< source
    accumulator tile
    TransformA transform_A = TransformA(),                ///<< transformation
    applied to A fragment
    TransformB transform_B = TransformB()) {              ///<< transformation
    applied to B fragment

    //
    // Prologue
    //

    // Perform accumulation in the 'd' output operand
    accum = src_accum;

    FragmentA tb_frag_A;
    FragmentB tb_frag_B;

    tb_frag_A.clear();
    tb_frag_B.clear();

    // The last kblock is loaded in the prolog
    iterator_A.load(tb_frag_A);
    iterator_B.load(tb_frag_B);

    ++iterator_A;
    ++iterator_B;

    this->smem_iterator_A.store(transform_A(tb_frag_A));
    this->smem_iterator_B.store(transform_B(tb_frag_B));

    ++this->smem_iterator_A;
    ++this->smem_iterator_B;

    __syncthreads();

    // Pair of fragments used to overlap shared memory loads and math
    instructions
    WarpFragmentA warp_frag_A[2];
    WarpFragmentB warp_frag_B[2];

    this->warp_tile_iterator_A.set_kgroup_index(0);
    this->warp_tile_iterator_B.set_kgroup_index(0);

    this->warp_tile_iterator_A.load(warp_frag_A[0]);
    this->warp_tile_iterator_B.load(warp_frag_B[0]);

    ++this->warp_tile_iterator_A;
    ++this->warp_tile_iterator_B;

```

```

Operator warp_mma;

int smem_write_stage_idx = 1;

// Avoid reading out of bounds
iterator_A.clear_mask(gemm_k_iterations <= 1);
iterator_B.clear_mask(gemm_k_iterations <= 1);

// Issue loads during the first warp-level matrix multiply-add *AFTER*
issuing
// shared memory loads (which have the tightest latency requirement).

//
// Mainloop
//

// Note: The main loop does not support Base::kWarpGemmIterations == 2.
CUTLASS_GEMM_LOOP
for (; gemm_k_iterations > 0; --gemm_k_iterations) {
    //
    // Loop over GEMM K dimension
    //

    CUTLASS_PRAGMA_UNROLL
    for (int warp_mma_k = 0; warp_mma_k < Base::kWarpGemmIterations;
++warp_mma_k) {

        // Load warp-level tiles from shared memory, wrapping to k offset if
this is the last group
        // as the case may be.

        if (warp_mma_k == Base::kWarpGemmIterations - 1) {

            // write fragments to shared memory
            this->smem_iterator_A.store(transform_A(tb_frag_A));

            this->smem_iterator_B.store(transform_B(tb_frag_B));

            __syncthreads();

            ++this->smem_iterator_A;
            ++this->smem_iterator_B;

            // Add negative offsets to return iterators to the 'start' of the
circular buffer in shared memory
            if (smem_write_stage_idx == 1) {
                this->smem_iterator_A.add_tile_offset({0, -Base::kStages});
                this->smem_iterator_B.add_tile_offset({-Base::kStages, 0});
            }
            else {
                this->warp_tile_iterator_A.add_tile_offset(
                    {0, -Base::kStages * Policy::kPartitionsK *
Base::kWarpGemmIterations});
                this->warp_tile_iterator_B.add_tile_offset(

```

```

        {-Base::kStages * Policy::kPartitionsK *
Base::kwarpGemmIterations,
        0});
    }

    smem_write_stage_idx ^= 1;
}

    this->warp_tile_iterator_A.set_kgroup_index((warp_mma_k + 1) %
Base::kwarpGemmIterations);
    this->warp_tile_iterator_B.set_kgroup_index((warp_mma_k + 1) %
Base::kwarpGemmIterations);

    this->warp_tile_iterator_A.load(warp_frag_A[(warp_mma_k + 1) % 2]);
    this->warp_tile_iterator_B.load(warp_frag_B[(warp_mma_k + 1) % 2]);

    ++this->warp_tile_iterator_A;
    ++this->warp_tile_iterator_B;

    if (warp_mma_k == 0) {

        iterator_A.load(tb_frag_A);
        iterator_B.load(tb_frag_B);

        ++iterator_A;
        ++iterator_B;


        // Avoid reading out of bounds if this was the last loop iteration
        iterator_A.clear_mask(gemm_k_iterations <= 2);
        iterator_B.clear_mask(gemm_k_iterations <= 2);
    }

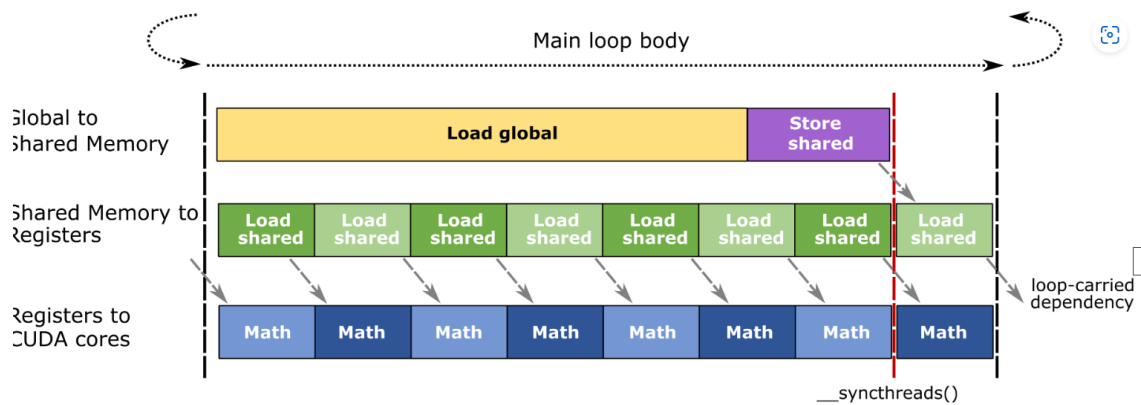
    warp_mma(accum, warp_frag_A[warp_mma_k % 2],
            warp_frag_B[warp_mma_k % 2], accum);
}
}

}

```

- analysis on code





## Reference

- [CUTLASS: Fast Linear Algebra in CUDA C++ | NVIDIA Technical Blog](#)
- [CUTLASS: Fast Linear Algebra in CUDA C++ - 知乎 \(zhihu.com\)](#) (translated version)
- [深入浅出GPU优化系列: GEMM优化 \(一\) - 知乎 \(zhihu.com\)](#)
- [深入浅出GPU优化系列: GEMM优化 \(二\) - 知乎 \(zhihu.com\)](#)
- [深入浅出GPU优化系列: GEMM优化 \(三\) - 知乎 \(zhihu.com\)](#)
- [SGEMM · NervanaSystems/maxas Wiki \(github.com\)](#)
- [CUDA 矩阵乘法终极优化指南 - 知乎 \(zhihu.com\)](#)
- [cuda\\_sgemm/gemm.cu at master · niuhope/cuda\\_sgemm \(github.com\)](#)
- [Why use CUTLASS instead of CUBLAS for GEMM? What are the advantages of CUTLASS ? · Issue #109 · NVIDIA/cutlass \(github.com\)](#)
- [Liu-xiandong/How to optimize in GPU: This is a series of GPU optimization topics. Here we will introduce how to optimize the CUDA kernel in detail. I will introduce several basic kernel optimizations, including: elementwise, reduce, sgemv, sgemm, etc. The performance of these kernels is basically at or near the theoretical limit. \(github.com\)](#)
- [\(27 封私信 / 80 条消息\) 自己写的CUDA矩阵乘法能优化到多快? - 知乎 \(zhihu.com\)](#)
- [CUDA矩阵乘法的优化 · wu-kan](#)
- [cutlass/efficient\\_gemm.md at master · NVIDIA/cutlass \(github.com\)](#)